

Fitting Model Predictive Control to an Epidemiological Model

Tamás Epres

July 2024

1 Introduction

The past few years have highlighted the fact that controlling and managing a pandemic is not a trivial task. Managing the progression of an epidemic is a complex problem that requires knowledge from multiple scientific disciplines; in this case, we may encounter issues that necessitate engineering problem-solving. In my work, I will primarily illuminate the engineering solutions and tasks necessary for controlling an epidemic from the perspective of control theory; these include the quantization of the control signal and the limitation of the duration of interventions. In solving this problem, I will apply model predictive control to various epidemiological models. Initially, I will introduce the applied strategy, the basics of model predictive control on a simple linear model, and then apply the strategy to a process model describing the epidemiological dynamics of a country. Last but not least, I aim to fit model predictive control to a multi-agent model mimicked by a neural network.

2 Introduction to Model Predictive Control

2.1 Definition of Model Predictive Control

In this section, I will focus primarily on the theoretical considerations of control, emerging problems, and not on the epidemic model. Model predictive control (hereafter referred to by the English abbreviation MPC) is a control strategy and method whose essence is to provide control sequences that minimize a certain error, knowing the dynamics that describe the process. In other words, MPC optimizes an error within a given time window and provides a control signal based on that. Let us start with a simple linear discrete-time process:

$$\mathbf{x}_{k+1} = \mathbf{A}\mathbf{x}_k + \mathbf{B}u_k \quad (1)$$

$$\mathbf{y}_k = \mathbf{C}\mathbf{x}_k \quad (2)$$

Where:

- \mathbf{y}_k is the response vector of the process at time k .
- \mathbf{x}_k is the state vector of the process at time k .
- u_k is the control input at time k .
- \mathbf{A} is the matrix describing the dynamics of the process, the system matrix.
- \mathbf{B} is the vector characteristic of the process that determines the control input.

For simplicity, let's deal with SISO systems and restrict our state variables to be one-dimensional. Thus, y_k and x_k are also one-dimensional, as is u_k . In addition, there is a cost function that defines the error that the algorithm seeks to minimize. In our first example, our goal is to bring the system from an initial state to another state prescribed by us. Thus, the optimization problem takes the following form:

$$\min \sum_{i=0}^N (\mathbf{x}_{k+i|k} - \mathbf{x}_{req})^2 \quad (3)$$

$$\mathbf{x}_{k+i+1|k} = \mathbf{A}\mathbf{x}_{k+i|k} + \mathbf{B}u_{k+i|k} \quad (4)$$

$$u_{down} < u_{k+i|k} < u_{up} \quad (5)$$

$$\mathbf{x}_{k+i|k} \in \mathbb{R}^n \quad (6)$$

$$u_{k+i|k} \in U \quad (7)$$

Where:

- N is the time horizon, up to which the system's responses are observed.
- $\mathbf{x}_{k+i|k}$ represents the responses of the process at time k along the time horizon with the variable i .
- $u_{k+i|k}$ is the control signal at time k along the time horizon with the variable i .
- u_{down} is the lower bound of the control signal.
- u_{up} is the upper bound of the control signal.
- \mathbf{x}_{req} is the desired state vector prescribed by us.
- $U \in \{U_0, U_1, \dots, U_9\}$

The objective of the optimization is to minimize equation (3), which is the objective, while equations (4) and (5) represent the constraints to be adhered to. It is evident that both u_k and x_k are the variables to be optimized. Furthermore, while the state vectors of the system can take continuous values, the control signals can only be selected from a set with finite cardinality. This optimization problem is known as a Mixed Integer Linear Problem (MILP).

2.2 Waiting Time Constraints (WTC)

As mentioned earlier, this section does not deal with the epidemiological model but rather addresses the issue of waiting time constraints. The basis of the problem is as follows: a given control action cannot be changed abruptly, and it is not advisable to maintain it for too long due to social and economic reasons. Therefore, we introduce constraints that regulate the duration of control actions, determining how long a control signal can remain active. In other words, we can apply a given control signal to the system at any time; however, if we intervene, the value of the intervention must be maintained for a minimum duration, and after a certain time, this value can no longer be applied until another control action is taken. To solve this, we introduce new variables to be optimized.

$$\mathbf{v}_j = [v_j^0, v_j^1, \dots, v_j^n]^T$$

Where:

$$v_j^l = \begin{cases} 1, & \text{if the given } u \text{ is active} \\ 0, & \text{otherwise} \end{cases}$$

We also assume that:

$$\sum_{l=0}^n v_j^l = 1 \quad (8)$$

Thus, \mathbf{v}_j is a vector where all elements are 0 except for the index corresponding to the value of u at that particular time. For example, if $u = U_1$ and $U \in \{U_0, U_1, \dots, U_9\}$, then $\mathbf{v}_j = [0, 1, 0, \dots, 0]$. We also introduce a counter \mathbf{H} that tracks how many v_j values have been active up to a given time point. If a v_j with a different value from the previous one appears in the sequence, the counter resets. Additionally, we introduce two parameters, \mathbf{W}_{Tmin} and \mathbf{W}_{Tmax} , which determine how long the given v_j index should be maintained and after what duration it should change its value. Their operation can be described as follows:

$$\mathbf{H}_{j+1} = (\mathbf{H}_j + \mathbf{v}_j) \odot \mathbf{v}_j \quad (9)$$

Where the \odot operator denotes element-wise multiplication, also known as the Hadamard product.

$$\mathbf{v}_j^T \cdot (\mathbf{W}_{Tmin} - \mathbf{H}_{j+1}) \leq M \cdot (\mathbf{v}_j^T \cdot \mathbf{v}_{j+1}) \quad (10)$$

Where M is a large number ensuring that the inequality holds when $\mathbf{v}_j = \mathbf{v}_{j+1}$.

$$\mathbf{H}_{j+1} \leq \mathbf{W}_{Tmax} \quad (11)$$

While equation (9) is responsible for counting, (10) establishes the lower bound, and (11) ensures the upper bound.

2.3 Simulation Results

In the following section, I present some simulation results. The optimization problem was implemented in Python using the Pyomo and NumPy packages, with visualization handled by the Matplotlib package. In this example, I used the Gurobi solver to solve the optimization problem since we are dealing with a convex problem. (A problem is convex if the objective function is convex, the equalities are affine, and the inequalities are convex.)

The parameters defining the system dynamics are:

$$\mathbf{A} = 1.1 \quad \mathbf{B} = 1.0$$

Additionally, the \mathbf{U} set, \mathbf{W}_{Tmin} , and \mathbf{W}_{Tmax} are defined as follows:

$$\mathbf{U} \in \{-10, -9, -8, \dots, 10\} \quad \mathbf{W}_{Tmin} = 2 \quad \mathbf{W}_{Tmax} = 4$$

Let the initial state \mathbf{x}_0 and the reference state \mathbf{x}_{req} be as follows:

$$\mathbf{x}_0 = 0.0 \quad \mathbf{x}_{req} = 100.$$




Figure 1: On the left, the control sequence; on the right, the evolution of the system

As can be seen, the system will oscillate around the response. This is because the control signal does not fit the dynamics and does not take on a value that would direct the system precisely to the reference signal. We can also observe that the oscillation consists of distinct segments, which is due to the restriction on the duration of interventions. It can be noted that if there were no constraints on the control signal, the system would reach the reference signal without oscillation.

For the sake of completeness, I extended the system to a two-dimensional case. In this case, the parameters were as follows:

$$\mathbf{A} = \begin{pmatrix} 1 & 1.4 \\ 1.4 & 1. \end{pmatrix} \quad \mathbf{B} = [1., 1.]$$

The \mathbf{U} set, as well as the \mathbf{W}_{Tmin} and \mathbf{W}_{Tmax} parameters, are identical to those in the previous one-dimensional case. The time horizon in this case, as in the one-dimensional case, was 30 s.

The initial state \mathbf{x}_0 and the reference state \mathbf{x}_{req} are as follows:

$$\mathbf{x}_0 = [0., 0.] \quad \mathbf{x}_{req} = [10., 10.]$$

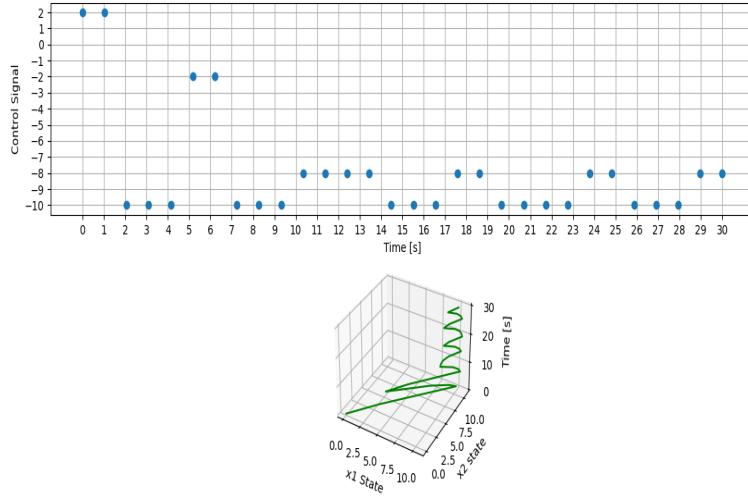


Figure 2: Above: control sequence; below: system evolution

3 Epidemic Process Model

3.1 Description of the Model Dynamics

In the following section, we will examine the epidemic model for which the strategy discussed in the previous chapter was developed, with slight modifications. The essence of the process model is as follows: the society is divided into different groups depending on the stage of infection they are in. The transition from one group to another is described by differential equations, meaning this

model is a continuous-time model. Let us see the individual groups and the description of the dynamics.

The uninfected population (S) consists of those who have not yet contracted the infection but may potentially become infected. The latent (L) population includes those who have encountered the infection but do not yet show any symptoms. Subsequently, the population moves to the pre-symptomatic infectious (P) group, from which individuals either progress to the symptomatic infected (I) group or the asymptomatic infected (A) group. Those who do not show symptoms are likely to recover from the disease, which requires defining the recovered (R) group. Those who show symptoms either recover and move to the R group or are hospitalized (H). Those in the hospital either recover or die (D). The dynamics are described by the following differential equations:

$$S'(t) = -\beta \frac{[P(t) + I(t) + \delta A(t)]S(t)}{N} \quad (12)$$

$$L'(t) = \beta \frac{[P(t) + I(t) + \delta A(t)]S(t)}{N} - \alpha L(t) \quad (13)$$

$$P'(t) = \alpha L(t) - pP(t) \quad (14)$$

$$I'(t) = qpP(t) - \rho_1 I(t) \quad (15)$$

$$A'(t) = (1 - q)pP(t) - \rho_A A(t) \quad (16)$$

$$H'(t) = \rho_1 \nu I(t) - hH(t) \quad (17)$$

$$R'(t) = \rho_1(1 - \nu)I(t) + \rho_A A(t) + (1 - \mu)hH(t) \quad (18)$$

$$D'(t) = \mu hH(t) \quad (19)$$

Where:

- β : The infection rate.
- δ : The infectiousness ratio in the $A(t)$ group.
- N : The total population size.
- α : The rate of transition from the $L(t)$ state.
- p : The rate of transition from the $P(t)$ state.
- q : The probability that an individual in the $P(t)$ state moves to the $I(t)$ state.
- ρ_1 : The recovery or transition rate of infected individuals ($I(t)$).
- ρ_A : The recovery or transition rate of asymptomatic individuals ($A(t)$).
- ν : The probability that an infected individual ($I(t)$) requires hospitalization.
- h : The mortality rate of hospitalized patients ($H(t)$).




Figure 3: Population dynamics over the course of the epidemic

The dynamics are illustrated in the following graph:

Since we need a discrete-time model for optimization, the system of differential equations mentioned above must be solved using a numerical method. In my solution, I used a fourth-order Runge-Kutta method. For numerical stability, I normalized the states by the size of the entire population.

Regarding the model, we can confirm that the model has a conserved property, meaning that if we sum all the equations, we would get exactly one when normalized by the population size. That is, no new individuals are created in the model. Furthermore, it can be established that, if the states are normalized in the system:

$$\lim_{t \rightarrow \infty} |S(t) + R(t) + H(t)| = 1$$

This means everyone will end up in either the S, R, or D groups. The model does not address the mechanism of reinfection.

We can now turn to the optimization problem. As shown in the previous chapter, the problem consists of an objective and its constraints.

$$\min \sum_{i=0}^N u_i^2 + (S_i - S_{term})^2 \quad (20)$$

$$\mathbf{x}_{i+1} = F(\mathbf{x}_i) \quad (21)$$

$$\mathbf{x}_i[5] \leq 20000 \quad (22)$$

Where:

- N is the time horizon.
- i is the running variable.
- u_i is the intervention (control input) introduced at time i .
- \mathbf{x}_i is the state vector of the system at time i .
- $\mathbf{x}_k[5]$ is the number of people in the hospital at time i .
- S_{term} is the terminal state of the system, i.e., the state of the S group when the epidemic has definitely subsided.

- S_i is the size of the S group at time i .

Equation (20) represents the objective of the pandemic, which is to minimize the value of interventions to reduce economic damage and bring the system to a state where the epidemic has subsided (\mathbf{x}_{term} terminal state). Equation (21) represents the sampling of the system dynamics, i.e., how we transition from the previous state to the next. Equation (22) imposes a constraint that the number of hospitalized patients should not exceed a certain value to prevent capacity issues in the public health system.

We can also address the concept of the terminal set or terminal state. A terminal state refers to the condition where, even with zero interventions, the number of hospitalized patients strictly decreases monotonically, indicating the development of herd immunity. Clearly, if we reach the terminal set, the epidemic will no longer spread. These are the states we aim to achieve during control, as reflected in the second term of the (20) expression, similar to how we aimed to achieve a reference state in the first example. Determining where these terminal states are is a very challenging problem. It is not trivial to determine whether a state is reachable in a nonlinear model. In my solution, I used a brute-force approach. By starting from random initial states, I observed the system to see whether those states were stable or not. The random initialization was implemented so that the individual groups would only take "reasonable" values and not unrealistic values.

3.2 Simulation Results

The solution was implemented in Python using the Pyomo and NumPy libraries, just as in the first chapter. The Baron solver was used to solve the optimization problem. Below are the parameters used for the simulation:

$$\mathbf{x}_{init} = [9,800,000-50, 50, 0., 0., 0., 0., 0., 0.]$$

That is, the total population is 9,800,000, while the number of people in the latent group is 50. Additionally, the time-related parameters are:

$$t_{end} = 180 \quad dt = 1$$

The total time horizon was set to 180 days, with the constraint that decisions can only be made weekly, and each decision must be maintained for one week. The sampling time is 1 day.

The following observations can be made: if we prescribe the number of people in the hospital, we can only reach the terminal set with a certain slope and speed, meaning time must pass to achieve herd immunity. Therefore, the larger the hospital capacity, the faster we can get through the epidemic. It is also evident that the size of the susceptible group S , and the groups capable of infection (L , P , A , I) is critical. Their relative sizes determine whether a second wave or a new peak in the hospital group H might occur with no further interventions.

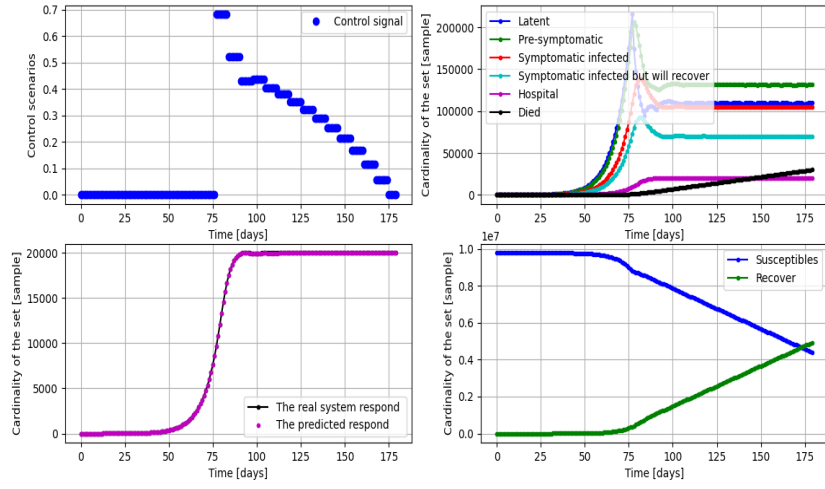


Figure 4: Changes in the system over time. The intervention sequence is shown in the upper left corner. The upper right corner shows the number of latent states and deceased individuals. The lower right corner displays the actual (continuous) and predicted (dashed) number of people in the hospital. The lower left corner shows the number of recovered and still uninfected individuals.